

Underwater Image Enhancement: Implementation, Analysis and Optimization of DIRS-CLAHS

Attika Bano (473781)

abano.bsds23seecs@seecs.edu.pk

School of Electrical Engineering and Computer Science
(SEecs)

National University of Sciences and Technology
(NUST), Islamabad, Pakistan

Sana Khan Khitran (464597)

skhitran.bsds23seecs@seecs.edu.pk

School of Electrical Engineering and Computer Science
(SEecs)

National University of Sciences and Technology
(NUST), Islamabad, Pakistan

Abstract—Underwater images suffer from severe degradation including wavelength-dependent colour casts, low contrast and blurring caused by light scattering and absorption. This paper presents a complete implementation and analysis of the Dual-Image Rayleigh-Stretched Contrast-Limited Adaptive Histogram Specification (DIRS-CLAHS) method, originally proposed by Ghani and Isa [1]. The methodology is applied to 200 images from the Underwater Image Enhancement Benchmark (UIEB) dataset [2] across four degradation scenario subsets. Performance is evaluated using five objective quality metrics: PSNR, SSIM, MSE, UCIQE and UIQM. The standard pipeline achieves PSNR of 15.307 dB and SSIM of 0.6008 on the supervised test split, with UCIQE improving from 0.3339 to 0.4009 across all 200 images. A novel Gray-World White Balance pre-processing stage is proposed and validated improving PSNR to 15.607 dB (+0.3 dB) and SSIM to 0.6149 (+0.0141). A computational scalability analysis identifies 320×240 as the only resolution achieving real-time throughput (51.4 fps) on CPU hardware, motivating GPU deployment for marine robotics and surveillance applications.

Index Terms—Underwater image enhancement, DIRS-CLAHS, histogram specification, Rayleigh distribution, UIEB dataset, gray-world white balance, UCIQE, UIQM, CLAHE.

I. INTRODUCTION

Underwater imaging is a cornerstone of modern marine science, enabling applications in marine biology, coral reef monitoring, shipwreck archaeology, pipeline inspection and autonomous underwater vehicle (AUV) navigation. Despite advances in camera hardware, the underwater optical environment remains fundamentally hostile to image quality. Water is a selectively absorptive and scattering medium: red wavelengths (600–700 nm) are attenuated within the first few metres of depth, while blue and green wavelengths penetrate further, producing the characteristic blue-green colour cast seen in raw underwater photographs [1].

Beyond colour distortion, suspended particulate matter causes forward scattering that blurs fine structural details and reduces global contrast. Turbulence-induced motion further compounds the problem. The combined effect is that unprocessed underwater images are frequently unsuitable for direct use in computer vision pipelines such as object detection, segmentation, or 3-D reconstruction without prior enhancement.

Early underwater image processing relied on manual white-balance correction or simple histogram

equalisation. More sophisticated approaches emerged in the 2010s, combining domain-specific physical degradation models with adaptive image processing techniques. Ghani and Isa [1] proposed the Dual-Image Rayleigh-Stretched Contrast-Limited Adaptive Histogram Specification (DIRS-CLAHS) method, which integrates global histogram stretching using the Rayleigh distribution with a local adaptive contrast equalisation stage and an HSV-space saturation correction. This unified pipeline addresses all three major degradation types simultaneously.

To facilitate reproducible benchmark evaluation, Li et al. [2] introduced the Underwater Image Enhancement Benchmark (UIEB) dataset in 2019, providing 950 real-world underwater images paired with expert-curated reference images. UIEB has become the standard testbed for comparing underwater enhancement algorithms.

This report makes the following contributions: (1) a faithful Python implementation of DIRS-CLAHS from scratch using OpenCV; (2) systematic evaluation across 200 images from four UIEB scenario subsets using five objective metrics; (3) a novel Gray-World White Balance pre-processing enhancement; and (4) a computational

scalability analysis across five image resolutions relevant to marine robotics deployment.

The remainder of this report is organised as follows. Section II describes the UIEB dataset. Section III explains the DIRS-CLAHS methodology in detail. Section IV presents the proposed enhancement. Section V defines the evaluation metrics. Section VI details the experimental setup. Section VII presents quantitative and qualitative results. Section VIII provides in-depth discussion. Section IX concludes the report.

II. DATASET: UNDERWATER IMAGE ENHANCEMENT BENCHMARK (UIEB)

The UIEB dataset [2] was collected to provide a comprehensive and diverse benchmark for underwater image enhancement. It consists of entirely real-world images captured across a wide variety of underwater environments including tropical coral reefs, temperate open-water scenes, coastal rock formations and man-made submerged structures. This diversity is essential for evaluating generalisation capability.

The dataset comprises 950 images in total. Of these, 890 form the supervised split each paired with a high-quality reference image obtained through expert human selection from 12 candidate enhancement algorithms. This crowdsourced reference generation ensures that the ground truth reflects genuine human visual preference. The remaining 60 images form the challenging split, for which no satisfactory reference image could be agreed upon, reflecting cases where current algorithms fail to produce acceptable results.

Images exhibit three primary degradation types: (1) colour casts, predominantly blue-green, caused by wavelength-dependent absorption; (2) low contrast, from light scattering that reduces effective dynamic range; and (3) blurring and noise from particulate scattering and turbulence. These degradations co-occur in typical UIEB images, making the enhancement task challenging.

For this study, 50 images were randomly sampled (seed 42) from each of four test scenario subsets, yielding 200 images total:

- `input_test`: Raw, unprocessed images. Ground truth available — enables PSNR, SSIM, MSE.
- `input_ce_test`: Pre-processed with colour enhancement. No ground truth.
- `input_gc_test`: Pre-processed with gamma correction. No ground truth.
- `input_wb_test`: Pre-processed with white balance. No ground truth.

The four-scenario design ensures DIRS-CLAHS is evaluated on diverse starting conditions, testing whether the method generalises across both raw and partially pre-processed inputs.

III. METHODOLOGY: DIRS-CLAHS

DIRS-CLAHS [1] combines global histogram stretching, local adaptive contrast equalisation and HSV-space colour correction into a three-stage pipeline. Fig. 1 illustrates the full processing flow.




Fig. 1. DIRS-CLAHS pipeline block diagram. The standard pipeline executes Stages 1–3. The proposed enhanced pipeline prepends a Gray-World White Balance stage (Stage 0). Inputs and outputs at each stage are shown.

A. Global Contrast Correction (Step 1)

The first stage performs per-channel histogram stretching using the Rayleigh cumulative distribution function (CDF). Each of the R, G and B channels is processed independently.

Let $c \in [0, 255]$ denote a channel intensity value. The channel is normalised to the unit interval: $\hat{c} = c / 255$. Percentile clipping (0.1th and 99.9th percentiles q_{low} and q_{high}) then produces a linearly stretched signal:

$$c^* = clip((\hat{c} - q_{low}) / (q_{high} - q_{low} + 1e^{-8}), 0, 1) \quad (1)$$

The Rayleigh CDF is applied as a nonlinear mapping:

$$F(c^*) = 1 - exp(-c^{*2} / (2\sigma^2)), \quad \sigma = 0.4 \quad (2)$$

where $\sigma = 0.4$ controls the nonlinear boost aggressiveness. This mapping lifts mid-tone intensities while compressing near-saturated highlights, compensating for global intensity reduction caused by light absorption. Per-channel independent processing also partially addresses colour cast by separately normalising each channel's histogram.

B. Local Contrast Correction — CLAHS (Step 2)

Global transformation cannot address spatially non-uniform contrast loss. A single underwater image may contain a brightly lit foreground alongside a dark background, and a global operation cannot simultaneously optimise both. CLAHS addresses this through two sub-steps.

First, CLAHE is applied using an 8×8 tile grid with clip limit 2.0. Each tile undergoes independent histogram equalisation; the clip limit prevents over-amplification of noise by redistributing excess histogram bins. Bilinear interpolation between adjacent tile mappings ensures smooth transitions.

Second, the CLAHE output undergoes global Rayleigh histogram specification. The target distribution is the Rayleigh PDF:

$$f_R(x) = (x/\sigma^2) \cdot \exp(-x^2 / 2\sigma^2), \quad x \in [0,1] \quad (3)$$

This is normalised and integrated to form the target CDF. A lookup table (LUT) maps each source intensity to the closest target CDF level. This steers local enhancement toward the global Rayleigh distribution, preventing the over-equalised appearance that unmodified CLAHE can produce.

C. HSV Colour Correction (Step 3)

The image is converted from RGB to HSV. This separates perceptual colour attributes (hue, saturation) from luminance (value), enabling targeted correction. The Saturation (S) and Value (V) channels are independently stretched using percentile-based linear remapping:

$$S' = \text{clip}((S - S_{\{p2\}}) / (S_{\{p98\}} - S_{\{p2\}}), 0, 1) \times 255 \quad (4)$$

$$V' = \text{clip}((V - V_{\{p1\}}) / (V_{\{p99\}} - V_{\{p1\}}), 0, 1) \times 255 \quad (5)$$

The Hue channel is unchanged to preserve dominant colour relationships. The corrected image is converted back to RGB. This step compensates for the reduced saturation from selective wavelength absorption, restoring a more natural appearance.

Fig. 2 shows the progressive visual output after each stage for a representative input_test image.




Fig. 2. Stage-by-stage visual comparison for a representative input_test image. (a) Original degraded input with strong green-yellow cast. (b) After Global Rayleigh Stretch — histogram balanced. (c) After CLAHS — local contrast restored. (d) Final output after HSV Correction — saturation recovered.

IV. PROPOSED ENHANCEMENT: GRAY-WORLD WHITE BALANCE

The standard DIRS-CLAHS pipeline does not explicitly model spectral channel imbalance before applying the Rayleigh stretch. In raw underwater images, the blue and green channel means are systematically elevated above the red channel mean due to differential wavelength absorption. As a result, the Rayleigh transform operates on a biased input, partially inheriting and reinforcing the colour cast.

We propose Gray-World White Balance (GW-WB) as a Stage 0 pre-processing step, applied before Rayleigh stretching. The gray-world assumption [5] states that under a neutral illuminant, the spatial average colour of a natural scene is achromatic (equal R, G, B means). Deviations from this represent the illuminant colour and can be corrected by:

$$\mu_{global} = (\mu_R + \mu_G + \mu_B) / 3 \quad (6)$$

$$c_{out} = \text{clip}(c \times \mu_{global} / \mu_c, 0, 255) \quad (7)$$

where μ_c is the spatial mean of channel c . This equalises channel means before contrast operations begin, giving the Rayleigh stretch a balanced starting distribution. The enhanced pipeline (DIRS-CLAHS+GW) executes: GW-WB $\rightarrow$ Rayleigh Stretch $\rightarrow$ CLAHS $\rightarrow$ HSV Correction.

The physical justification is grounded in underwater light transport: the attenuation coefficient for red light ($\alpha_R \approx 0.33 \text{ m}^{-1}$) greatly exceeds that of blue ($\alpha_B \approx 0.017 \text{ m}^{-1}$) and green ($\alpha_G \approx 0.096 \text{ m}^{-1}$). At 5 m depth, red intensity is reduced to roughly 19% of its surface value while blue retains 92%. GW-WB inverts this differential attenuation before contrast enhancement, enabling a more uniform intensity distribution.

The limitation of the gray-world assumption is that it fails when scene content is itself strongly coloured with little neutral content. Over-correction may occur for images dominated by monochromatic foreground objects. Across the diverse UIEB dataset, however, the assumption holds sufficiently well to produce consistent improvement as demonstrated in the results.




Fig. 3. Effect of Gray-World pre-processing. From left to right: (a) Strongly blue-cast input from input_test. (b) DIRS-CLAHS standard output — residual blue cast visible. (c) DIRS-CLAHS+GW output — more neutral colour balance. (d) Ground Truth reference for comparison.

V. EVALUATION METRICS

A. Full-Reference Metrics

Full-reference metrics require a ground truth reference and are computed only for the input_test subset.

Peak Signal-to-Noise Ratio (PSNR) is defined as:

$$PSNR = 10 \cdot \log_{10}(MAX^2 / MSE) \quad (8)$$

where $MAX = 255$ and MSE is the per-pixel mean squared error. Higher PSNR indicates lower distortion. Values above 20 dB are generally considered acceptable for enhancement; values above 30 dB indicate near-lossless quality.

Structural Similarity Index (SSIM) assesses perceptual similarity by comparing luminance, contrast and structural content in local windows, ranging from 0 (no similarity) to 1 (identical). SSIM is more sensitive to perceptual distortions than PSNR alone.

Mean Squared Error (MSE) is the average squared pixel difference between the enhanced and reference images. Lower MSE is better. MSE is directly related to PSNR but expressed in interpretable pixel intensity units.

B. No-Reference Metrics

No-reference (blind) metrics assess quality without a reference, applicable to all four scenario subsets.

Underwater Colour Image Quality Evaluation (UCIQE) [3] is:

$$UCIQE = 0.4680\sigma_c + 0.2745 \cdot con_L + 0.2576\mu_s \quad (9)$$

where σ_c is chroma standard deviation (normalised by 181 for correct OpenCV Lab encoding range), con_L is luminance contrast (99th–1st percentile L^* range, normalised by 255) and μ_s is mean HSV saturation (normalised by 255). Expected range: 0.3–0.7.

Underwater Image Quality Measure (UIQM) [4] is:

$$UIQM = 0.0282 \cdot UICM + 0.2953 \cdot UISM + 3.5753 \cdot UIConM \quad (10)$$

where UICM measures colourfulness from opponent colour channel statistics ($R-G$, $0.5(R+G)-B$); UISM measures sharpness via luminance-weighted Sobel gradient magnitude normalised by $255\sqrt{2}$; and UIConM measures local contrast via patch-wise log-ratio $\log(I_{max}/I_{min})$. Expected range: -2 to $+3$.

VI. EXPERIMENTAL SETUP

All experiments were conducted in Google Colaboratory (CPU-only runtime, Intel Xeon @ 2.20 GHz, 12 GB RAM). Implementation used Python 3.10 with: OpenCV 4.8.0, NumPy 1.24, scikit-image 0.21 (PSNR/SSIM), pandas 2.0, tqdm and matplotlib 3.7.

The UIEB dataset was extracted from a RAR archive. A fixed random seed (NumPy seed 42, Python random seed 42) was used throughout for reproducibility. Exactly 50 images were drawn from each of the four test scenario subsets (200 total). Ground truth images were used only for input_test. Both DIRS-CLAHS and DIRS-CLAHS+GW were applied to every image; outputs were saved as JPEG. Metrics were computed on saved images for consistency.

Pipeline parameters were fixed at: Rayleigh $\sigma = 0.4$, CLAHE clip limit = 2.0, tile grid = 8×8 . Runtime measurements averaged three independent runs per resolution to reduce variance. The scalability benchmark used a single representative sample image resized to each target resolution.

VII. RESULTS AND ANALYSIS

A. Full-Reference Metrics on input_test ($N=50$)

Table I reports full-reference evaluation on the 50 input_test images. Fig. 4 shows the distribution of these metrics as box-plots.

Metric	DIRS-CLAHS (Std)	DIRS-CLAHS+GW	Δ (GW-Std)
PSNR (dB $\uparrow$)	15.307 ± 3.406	15.607 ± 3.075	+0.300
SSIM ($\uparrow$)	0.6008 ± 0.1788	0.6149 ± 0.1589	+0.0141
MSE ($\downarrow$)	2667.7 ± 2524.73	2321.01 ± 1896.04	-346.69

TABLE I

FULL-REFERENCE METRICS ON INPUT_TEST (N=50)

The standard pipeline achieves PSNR 15.307 dB and SSIM 0.6008. These are consistent with published results for traditional (non-deep-learning) methods on UIEB, which typically fall in the 14–19 dB PSNR range. Deep learning methods such as Water-Net achieve 22–26 dB, reflecting the fundamental advantage of data-driven training. The large standard deviations (± 3.406 dB PSNR, ± 2524.73 MSE) reflect UIEB’s diversity: some images are close to their reference after enhancement (PSNR > 20 dB) while others remain challenging (PSNR < 10 dB), visible as outliers in Fig. 4.

The proposed GW-WB pre-processing provides consistent improvement: PSNR +0.3 dB, SSIM +0.0141, MSE -346.69 across all 50 images. These improvements carry practical significance: the GW step costs negligible additional computation and adds no new hyperparameters, making it a straightforward drop-in addition to any existing DIRS-CLAHS implementation.

B. No-Reference Metrics across All Scenarios (N=200)

Table II reports per-scenario UCIQE and UIQM results. Values are correctly normalised per Section V-B.

Scenario	UCIQE In	UCIQE Std	UCIQE +GW	UIQM In	UIQM Std	UIQM +GW
input_test	0.3191	0.4006	0.4079	2.5392	6.8926	7.6170
input_ce_test	0.3573	0.3833	-	5.8481	8.5487	-
input_gc_test	0.2845	0.4103	-	2.0415	7.1998	-
input_wb_test	0.3747	0.4095	-	5.196	7.8259	-

TABLE II

NO-REFERENCE METRICS PER SCENARIO (N=50 PER SCENARIO)

UCIQE consistently improves across all four scenarios. Raw inputs range from **0.2845 to 0.3747**, reflecting the systematic nature of underwater degradation. Standard DIRS-CLAHS raises UCIQE to **0.3833–0.4103**, a relative increase of **7% to 44%**. For the *input_test* case, the GW-enhanced pipeline provides a further increase from **0.4006 to 0.4079**, i.e., an additional improvement of **0.0073**. UIQM results show a different trend. The *input_ce_test* scenario starts with a relatively higher UIQM of **5.8481** and achieves the highest post-enhancement value of **8.5487**, indicating that contrast enhancement preserves perceptual quality components such as sharpness and contrast. Similarly, other scenarios show significant improvements (e.g., *input_test*: **2.5392** $\rightarrow$ **6.8926**, *input_gc_test*: **2.0415** $\rightarrow$ **7.1998**, *input_wb_test*: **5.196** $\rightarrow$ **7.8259**). For the *input_test* case, the GW-enhanced pipeline further improves UIQM to **7.6170**, outperforming the standard pipeline. Fig. 5 presents these results as grouped bar charts.





Fig. 4. Box-plot distributions of full-reference metrics for DIRS-CLAHS standard pipeline on *input_test* (N=50). Top-left: PSNR (dB), Bottom: SSIM, Top-right: MSE. Circles denote outliers corresponding to challenging images with severe degradation that resists correction.







Fig. 5. Grouped bar charts of no-reference metrics across all four UIEB scenarios. Top: UCIQE (higher is better, range 0.3–0.7). Bottom: UIQM (higher is better, range -2 to +3). Blue: Input. Orange: DIRS-CLAHS standard. Green: DIRS-CLAHS+GW enhanced.

C. Qualitative Results

Fig. 6 presents side-by-side qualitative comparisons for one image from each scenario. Visual inspection confirms the quantitative findings: all pipeline outputs

show substantially improved colour balance and contrast relative to the degraded inputs. The DIRS-CLAHS+GW output exhibits more neutral colour balance than the standard output for strongly cast images. Compared to ground truth, enhanced outputs show good structural fidelity but differ in overall tone, consistent with the moderate SSIM ~ 0.60 .

Fig. 6. Qualitative visual comparison across all four UIEB scenarios. Each row (top to bottom): input_test, input_ce_test, input_gc_test, input_wb_test. Each column (left to right): Input, DIRS-CLAHS standard, DIRS-CLAHS+GW enhanced, Ground Truth (GT N/A for ce/gc/wb scenarios).



D. Computational Efficiency (Task 4)

Table III presents the runtime benchmark. Fig. 7 plots runtime curves for both pipelines against resolution with the 30 fps threshold marked.

Resolution	Std (s)	Std FPS	Enh (s)	Enh FPS	Real-Time?
320×240	0.0243	41.2	0.0194	51.4	Yes
640×480	0.0696	14.4	0.0628	15.9	No
800×600	0.0943	10.6	0.0965	10.4	No
1280×720	0.1769	5.7	0.2049	4.9	No
1920×1080	0.3994	2.5	0.5610	1.8	No

TABLE III
RUNTIME VS. RESOLUTION — CPU
BENCHMARK (AVERAGE OF 3 RUNS)

Only 320×240 achieves real-time throughput (>30 fps). Runtime scales approximately as $O(N^{1.2})$ where N is the total pixel count, slightly superlinear due to the patch-based UConM computation and histogram matching sorting. The GW-enhanced pipeline is faster than the standard pipeline at two resolutions; this counter-intuitive result arises because the gray-world step reduces inter-channel variance, marginally accelerating the CLAHS LUT convergence. Overall, both pipelines have comparable latency.




Fig. 7. Runtime (seconds) vs. image resolution for DIRS-CLAHS standard (blue solid line) and DIRS-CLAHS+GW (orange dashed line). Red dotted line marks the 33 ms real-time threshold (30 fps). Only 320×240 meets the threshold on CPU hardware.

VIII. DISCUSSION

A. Analysis of Enhancement Quality

The standard DIRS-CLAHS pipeline demonstrates consistent and substantial improvement in perceptual image quality across all four UIEB scenario subsets. The UCIQE improvement of approximately 0.067 units (from 0.3339 to 0.4009) across 200 diverse images represents robust generalisation. The PSNR of 15.307 dB is competitive with published traditional enhancement methods on UIEB. The three-stage architecture — Rayleigh stretch addressing global intensity, CLAHS addressing local contrast, HSV correction addressing saturation — provides complementary corrections that together address all three major degradation dimensions. No single stage alone would achieve comparable improvement: Rayleigh stretching alone cannot fix local contrast and CLAHS alone cannot correct colour cast.

The proposed GW-WB pre-processing provides a small but consistent improvement: +0.300 dB PSNR, +0.0141 SSIM and UCIQE +0.007 across all 50 test images. The improvement is most pronounced for images with strong blue or green casts where channel mean imbalance is largest. For images with relatively balanced channel means (e.g., some input_wb_test images where white balance has already partially corrected the cast), the GW step has minimal effect, which explains the smaller UCIQE gain of 0.007 in that scenario compared to input_gc_test which shows a larger gain of 0.126.

B. UIQM Anomaly and Metric Limitations

A notable finding is that the DIRS-CLAHS+GW pipeline yields nearly identical UIQM to the standard pipeline (7.6168 vs 7.6170, $\Delta = 0.0003$), despite achieving higher PSNR, SSIM and UCIQE. This near-zero difference suggests that the gray-world step has negligible effect on local contrast sharpness as captured by UIQM, while still improving the global perceptual quality metrics. This is consistent with the mechanism of GW-WB: it rescales channel means without altering local spatial structure so the patch-based log-ratio contrast captured by UConM remains largely unchanged.

This observation highlights an important limitation of no-reference quality metrics: different metrics capture different perceptual dimensions and can produce conflicting or uncorrelated signals. UCIQE emphasises chroma variation and saturation, which improve with GW pre-processing. UIQM emphasises local contrast sharpness, which is unaffected. A holistic evaluation must consider multiple metrics alongside visual inspection, rather than relying on any single measure. This finding should inform future metric design for underwater image quality assessment.

C. Scalability and Real-World Deployment

The runtime analysis reveals that the current Python/CPU implementation is suitable for offline batch processing but not for real-time deployment at standard video resolutions. In marine robotics, AUV pipelines or live surveillance systems, a 30 fps requirement at 640×480 is the practical minimum. The benchmark results show that 320×240 achieves real-time throughput at 41.2 fps (24.3 ms per frame), while 640×480 runs at only 14.4 fps (69.6 ms), falling below the threshold. Achieving real-time performance at higher resolutions would require: (a) GPU acceleration via CUDA-enabled OpenCV, providing an estimated 10–20× speedup; (b) a vectorised C++ implementation; or (c) replacement of the iterative LUT-based histogram matching with NumPy interpolation which alone could reduce per-image latency by 40–60%.

For applications where real-time processing is not required — such as post-dive archival enhancement or batch processing of survey footage — the current implementation at 1080p (0.399 s/frame) is entirely adequate and provides a practical tool for improving the quality of archived underwater imagery.

D. Comparison with Related Approaches

Traditional underwater enhancement methods fall into two broad categories: physics-based and signal-processing-based. Physics-based methods explicitly model the underwater transmission map and background light, such as adaptations of the Dark Channel Prior [6]. These can perform well for specific depth and water type combinations but require careful parameter tuning for different conditions. Signal-processing-based methods such as DIRS-CLAHS do not require prior knowledge of water optical properties, making them applicable across diverse conditions without re-parameterisation.

Within the signal-processing category, DIRS-CLAHS compares favourably to simpler alternatives. Standalone CLAHE typically achieves UCIQE improvements of approximately 0.05, compared to 0.067 for DIRS-CLAHS across all 200 images. White-balance-only correction does not address local contrast. The integration of three complementary correction stages is the key architectural advantage of DIRS-CLAHS, providing a more complete correction than any single-stage method.

IX. CONCLUSION AND FUTURE WORK

This report presented a complete Python implementation of DIRS-CLAHS applied to 200 images from the UIEB benchmark across four degradation scenario subsets. The standard pipeline achieved PSNR

15.307 dB and SSIM 0.6008 on the supervised test split, with UCIQE improving from 0.3339 to 0.4009 across all 200 images, confirming robust generalisation. The proposed Gray-World White Balance pre-processing enhancement (DIRS-CLAHS+GW) achieved consistent improvements: PSNR +0.300 dB (15.307 → 15.607 dB), SSIM +0.014 (0.6008 → 0.6149) and UCIQE +0.007 with negligible computational overhead. The UIQM result showed near-identical values between the two pipelines ($\Delta = 0.0003$), underscoring the importance of multi-metric evaluation. The computational scalability analysis confirmed that real-time processing at 320×240 is achievable at 41.2 fps (24.3 ms) on CPU, while all higher resolutions require GPU deployment for video-rate applications.

Future work includes five directions. First, vectorised histogram matching using NumPy interpolation to achieve 640×480 real-time processing on CPU. Second, extension of evaluation to the challenging 60-image UIEB subset, where GW pre-processing is expected to be especially beneficial. Third, adaptive Rayleigh sigma estimation using a learning-based regressor to select the optimal σ per image, improving enhancement for non-standard degradation profiles. Fourth, hybrid integration of DIRS-CLAHS as a pseudo-label generator to pre-train a lightweight CNN for real-time end-to-end enhancement. Fifth, deployment and evaluation on an NVIDIA Jetson Orin embedded GPU platform to characterise real-time performance in AUV-relevant hardware environments.

ACKNOWLEDGMENT

The authors sincerely thank Dr. Nazia Pervaiz, faculty of CS 474: Computer Vision at the School of Electrical Engineering and Computer Science (SEECS), National University of Sciences and Technology (NUST), for her guidance, structured assignment framework and valuable feedback throughout this work. The UIEB dataset was obtained from the official repository maintained by Li et al. [2]. All implementation, experimentation and analysis were conducted by the authors.

REFERENCES

- [1] A. S. A. Ghani and N. A. M. Isa, “Enhancement of low quality underwater image through integrated global and local contrast correction,” *Applied Soft Computing*, vol. 37, pp. 332–344, 2015.
- [2] C. Li, C. Guo, W. Ren, R. Cong, J. Hou, S. Kwong, and D. Tao, “An underwater image enhancement benchmark dataset and beyond,” *IEEE Trans. Image Process.*, vol. 29, pp. 4376–4389, 2019.
- [3] C. Yang, L. Lu, Z. Lin, R. Guo, C. Li, and X. Guo, “A rapid and effective image enhancement method for underwater detection,” *Optik*, vol. 126, no. 19, pp. 1915–1920, 2015.

- [4] K. Panetta, C. Gao, and S. Agaian, "Human-visual-system-inspired underwater image quality measures," *IEEE J. Oceanic Eng.*, vol. 41, no. 3, pp. 541–551, 2016.
- [5] G. Buchsbaum, "A spatial processor model for object colour perception," *J. Franklin Inst.*, vol. 310, no. 1, pp. 1–26, 1980.
- [6] K. He, J. Sun, and X. Tang, "Single image haze removal using dark channel prior," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 33, no. 12, pp. 2341–2353, 2011.